

# Lecture 8: Advanced topics

Hado van Hasselt

UCL 2018

# Overview of the course

1. **Learning to make decisions** in bandit problems; exploration vs exploitation; learning action values; greedy and  $\epsilon$ -greedy; policy gradient for bandits; UCB
2. **Sequential decision problems**; MDPs; planning with dynamic programming; policy evaluation + policy improvement = policy iteration
3. **Model-free prediction and control**; Monte Carlo returns; TD learning; on-policy; off-policy; Q-learning; Sarsa; Double Q-learning
4. **Function approximation and deep RL**; tabular vs linear vs non-linear; convergence and divergence; least-squares prediction (LSTD and LSMC); multi-step returns; neural Q-learning; DQN
5. **Policy gradients and actor-critic methods**; REINFORCE; advantage actor-critics (A2C); trust-region methods; continuous actions; CACLA; gradient ascent on the value (DPG)
6. **Learning from a model**; Full models vs expectation models vs stochastic (generative) models; Dyna; parametric vs non-parametric models; experience replay; search; MCTS

## Advanced topics and active research

- ▶ The main question is: **how do we maximize future rewards**
- ▶ Some main sub-questions are:
  - ▶ What do we learn? (Predictions, models, policies, ...)
  - ▶ How do we learn it? (TD, planning, ...)
  - ▶ How do we represent the learnt knowledge? (deep networks, sample buffers, ...)
  - ▶ How do we use the learnt knowledge?
- ▶ Specific active research topics include:
  - ▶ **Exploration** in the full sequential, function approximation case
  - ▶ **Credit assignment** with very delayed rewards
  - ▶ **Planning** with partial or **inaccurate models**
  - ▶ **Sample efficient** learning
  - ▶ **Appropriate generalization** (e.g., fast learning in new situations)
  - ▶ Building a useful, general, and information-rich **agent state**

## Case study: rainbow DQN (Hessel et al. 2018)

- ▶ Investigation of several algorithm components
- ▶ The starting point was DQN, with **target networks** and **experience replay**
- ▶ The components were:
  - ▶ **Double Q-learning**
  - ▶ **Prioritized replay**
  - ▶ **Splitting values from advantages** ('dueling network architectures')
  - ▶ **Multi-step updates**
  - ▶ **Distributional reinforcement learning**
  - ▶ **Parameter noise for exploration** ('noisy networks')
- ▶ We combined all components, and looked at performance

## Domain: Arcade Learning Environment (Bellemare et al. 2013)

- ▶ We use Atari games from the ALE as benchmark
  - ▶ Diverse set of games
  - ▶ Fun and interesting for humans
  - ▶ Good level of difficulty to test algorithms
  - ▶ Simulation is easy to work with — good for testing ideas
- ▶ The goal is to build a **general learning algorithm** without game-specific knowledge
- ▶ We will allow some Atari-specific knowledge (e.g., size of input screen)
- ▶ Can we build an agent that can play all (or most) games well?

## Starting point: DQN (Mnih et al. 2013, 2015)

- ▶ DQN includes:

- ▶ Convolutional neural network  $q_\theta: O_t \mapsto \mathbb{R}^m$  for  $m$  actions
- ▶  $\epsilon$ -greedy policy:  $\pi_t$
- ▶ Replay buffer for experience replay
- ▶ Target network parameters  $\theta^-$  (initially  $\theta_0^- = \theta_0$ )
- ▶ Q-learning loss function on  $\theta$  (uses replay and target network)

$$l(\theta) = \frac{1}{2} \left( R_{i+1} + \gamma \mathbb{E}_a \max_a q_{\theta^-}(S_{i+1}, a) - q_\theta(S_i, A_i) \right)^2$$

- ▶ Optimization method to minimize the loss (e.g., SGD, RMSprop, or Adam)
- ▶ Update  $\theta_t^- \leftarrow \theta_t$  occasionally  
(e.g., every 10000 steps — on all other steps  $\theta_t^- = \theta_{t-1}^-$ )

## Double DQN (van Hasselt et al. 2016)

$$l(\theta) = \frac{1}{2} \left( R_{i+1} + \gamma \llbracket q_{\theta}(S_{i+1}, \underset{a}{\operatorname{argmax}} q_{\theta}(S_{i+1}, a)) \rrbracket - q_{\theta}(S_i, A_i) \right)^2$$

## Prioritized replay (Schaul et al. 2016)

- ▶ DQN samples uniformly from replay
- ▶ Idea: prioritize transitions on which we can learn much
- ▶ Basic implementation:

$$\text{priority of sample } i = |\delta_i|,$$


where  $\delta_i$  was the TD error on the last this transition was sampled

- ▶ Sample according to priority
- ▶ Typically involves some additional design choices



## Dueling networks (Wang et al. 2016)

- ▶ We can decompose  $q_{\theta}(s, a) = v_{\xi}(s) + A_{\chi}(s, a)$ , where  $\theta = \xi \cup \chi$
- ▶ Here  $A_{\chi}(s, a)$  is the **advantage** for taking action  $a$

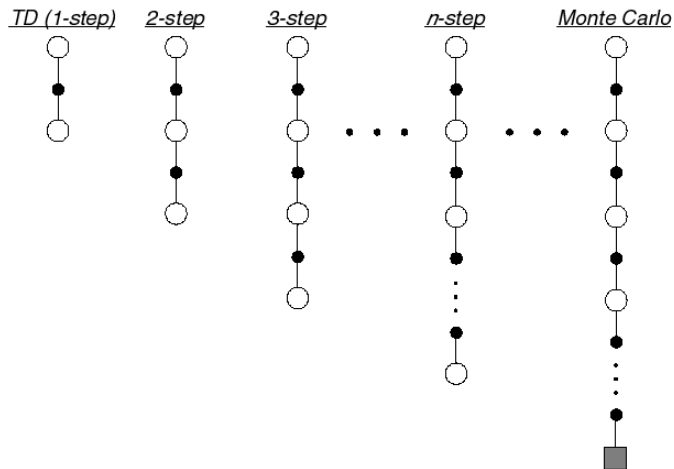


## Dueling networks

Video

## \yellow{Multi-step updates} (Sutton 1988)

- Let TD target look  $n$  steps into the future



## \yellow{Multi-step updates}

- ▶ Consider the following  $n$ -step returns for  $n = 1, 2, \infty$ :

$$n = 1 \quad (TD) \quad G_t^{(1)} = R_{t+1} + \gamma v(S_{t+1})$$

$$n = 2 \quad G_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 v(S_{t+2})$$

$$\vdots$$

$$n = \infty \quad (MC) \quad G_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t-1} R_T$$

- ▶ Define the  $n$ -step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n v(S_{t+n})$$

- ▶  $n$ -step temporal-difference learning

$$v(S_t) \leftarrow v(S_t) + \alpha \left( G_t^{(n)} - v(S_t) \right)$$

## \yellow{Multi-step updates}

- ▶ Define the  $n$ -step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \underbrace{\gamma^n q_{\theta}(S_{t+n}, \underset{a}{\operatorname{argmax}} q_{\theta}(S_{t+n}, a))}_{\text{Double Q bootstrap target}}$$

- ▶ Multi-step Q-learning

$$q(S_t, A_t) \leftarrow q(S_t, A_t) + \alpha \left( G_t^{(n)} - q(S_t, A_t) \right)$$

- ▶ Return is partially on-policy, bootstrap is off-policy
- ▶ That's okay — less greedy, but still policy improvement
- ▶ Still a well-defined prediction target:  
*“what if I'm on-policy for  $n$  steps, and then take a greedy action?”*

## Distributional reinforcement learning (Bellemare, Dabney, Munos 2017)

- ▶ So far, we've focused on learning **expected cumulative rewards**
- ▶ We can consider learning other things
- ▶ One possibility: learning the **distribution of returns**
- ▶ Knowing the distribution may be helpful for some things
- ▶ For instance, you can perhaps reason about the probability of termination
- ▶ It also means our representation (e.g., deep neural network) is forced to learn more
- ▶ This can speed up learning: learning more means potentially fewer samples

# Distributional reinforcement learning

- ▶ A specific instance is **Categorical DQN** (Bellemare et al., 2017)
- ▶ Consider a ‘comb’ distribution on  $\mathbf{z} = (-10, -9.9, \dots, 9.9, 10)^\top$
- ▶ For each point of support, we assign a ‘probability’  $p_\theta^i(S_t, A_t)$
- ▶ The approximate distribution of the return  $s$  and  $a$  is the tuple  $(\mathbf{z}, \mathbf{p}_\theta(s, a))$
- ▶ Our estimate of the expectation is:  $\mathbf{z}^\top \mathbf{p}_\theta(s, a) \approx q(s, a)$  — use this to act
- ▶ Goal: learn these probabilities

# Distributional reinforcement learning

1. Find max action:

$$a^* = \underset{a}{\operatorname{argmax}} \mathbf{z}^\top \mathbf{p}_\theta(S_{t+1}, a)$$

(use, e.g.,  $\theta^-$  for double Q)

2. Update support:

$$\mathbf{z}' = R_{t+1} + \gamma \mathbf{z}$$


3. Project distribution  $(\mathbf{z}', \mathbf{p}_\theta(S_{t+1}, a^*))$  onto support  $\mathbf{z}$

$$d' = (\mathbf{z}, \mathbf{p}') = \Pi(\mathbf{z}', \mathbf{p}_\theta(S_{t+1}, a^*))$$

where  $\Pi$  denotes projection

4. Minimize divergence

$$\text{KL}(d' \| d) = - \sum_i p'_i \frac{\log p'_i}{\log p_\theta^i(S_t, A_t)}$$



Bottom-right: target distribution

$$\Pi(R_{t+1} + \gamma \mathbf{z}, \mathbf{p}_\theta(S_{t+1}, a^*))$$

Update  $\mathbf{p}_\theta(S_t, A_t)$  towards this



## Noisy networks (Fortunato, Azar, Piot, et al. 2017)


- ▶ DQN used  $\epsilon$ -greedy exploration
- ▶ We learnt that UCB is better in bandits, but this is hard with function approximation
- ▶ Idea: add noise to parameters, replace all linear operations  $\mathbf{y} = \mathbf{W}\mathbf{x} + \mathbf{b}$  with

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \mathbf{b} + (\mathbf{W}' \times \varepsilon^{\mathbf{W}})\mathbf{x} + \mathbf{b}' \times \varepsilon^{\mathbf{b}}$$


where ' $\times$ ' is element-wise product and  $\varepsilon^{\mathbf{W}}$  and  $\varepsilon^{\mathbf{b}}$  are random matrix and vector

- ▶ The algorithm can learn to set  $\mathbf{W}' = \mathbf{0}$  and  $\mathbf{b}' = \mathbf{0}$ , but that takes time
- ▶ In the meantime, output is stochastic, and more so for rarely seen inputs  $\mathbf{x}$
- ▶ Results in exploration

## Rainbow DQN: results



## Rainbow DQN: ablation results



## Rainbow DQN: conclusions

- ▶ Components work well together
- ▶ Most important: prioritising replay, multi-step returns
- ▶ Least important: double, dueling
- ▶ No wild overestimations due to fixed bounded support of value distribution
- ▶ But this requires knowing appropriate range...
- ▶ ...but different game have different score ranges
- ▶ This is possible due to **reward clipping**: in DQN rewards are clipped to  $[-1, 1]$
- ▶ Makes learning easier, but changes the objective...

## Adaptive target normalization (van Hasselt et al. 2016)

- ▶ We can normalize targets before doing an update
- ▶ In online RL, we do not have access to the full 'data-set'
- ▶ In fact, data will change when our policy changes
- ▶ Solution: **adaptive normalization of updates**

## Adaptive target normalization (van Hasselt et al. 2016)

1. Observe target, e.g.,  $T_{t+1} = R_{t+1} + \gamma \max_a q_\theta(S_{t+1}, a)$
2. Update normalization parameters:

$$\mu_{t+1} = \mu_t + \eta(T_{t+1} - \mu_t) \quad \text{(first moment / mean)}$$

$$\nu_{t+1} = \nu_t + \eta(T_{t+1}^2 - \nu_t) \quad \text{(second moment)}$$

$$\sigma_{t+1} = \nu_{t+1} - \mu_{t+1}^2 \quad \text{(variance)}$$

where  $\eta$  is a step size (e.g.,  $\eta = 0.001$ )

3. Network outputs  $\tilde{q}_\theta(s, a)$ , update with

$$\Delta\theta_t \propto \left( \frac{T_{t+1} - \mu_{t+1}}{\sigma_{t+1}} - \tilde{q}_\theta(S_t, A_t) \right) \nabla_\theta \tilde{q}_\theta(S_t, A_t)$$

4. Recover **unnormalized** value:  $q_\theta(s, a) = \sigma_t \tilde{q}_\theta(s, a) + \mu_t$  (used for bootstrapping)

## Preserve outputs

- ▶ Naive implementation changes **all outputs** whenever we update the normalization
- ▶ This seems bad: we should avoid updating values of unrelated states
- ▶ We can avoid this. Typically:

$$\tilde{\mathbf{q}}_{\mathbf{W}, \mathbf{b}, \theta}(s) = \mathbf{W} \phi_{\theta}(s) + \mathbf{b}.$$

- ▶ Idea: define

$$\mathbf{W}'_t = \frac{\sigma_t}{\sigma_{t+1}} \mathbf{W}$$

$$\mathbf{b}'_t = \frac{\sigma_t \mathbf{b}_t + \mu_t - \mu_{t+1}}{\sigma_{t+1}}$$



Then

$$\sigma_{t+1} \tilde{\mathbf{q}}_{\mathbf{W}'_t, \mathbf{b}'_t, \theta_t}(s) + \mu_{t+1} = \sigma_t \tilde{\mathbf{q}}_{\mathbf{W}_t, \mathbf{b}_t, \theta_t}(s) + \mu_t$$

- ▶ Then update  $\mathbf{W}'_t$ ,  $\mathbf{b}'_t$  and  $\theta_t$  as normal (e.g., SGD)

## Preserve outputs

- Preserve Outputs Precisely, while Adaptively Rescaling Targets: Pop-Art






Pop-Art

Video

## Obtaining knowledge

- ▶ We want our agents to understand and interact with their environment
- ▶ A single reward signal can be sparse, and low in information
- ▶ But we can use the same prediction methods to learn **many things**



Horde (Sutton et al. 2011)

# General value functions

- ▶ Key idea is to consider general value functions (GVF)
- ▶ A GVF is conditioned on more than just state and actions

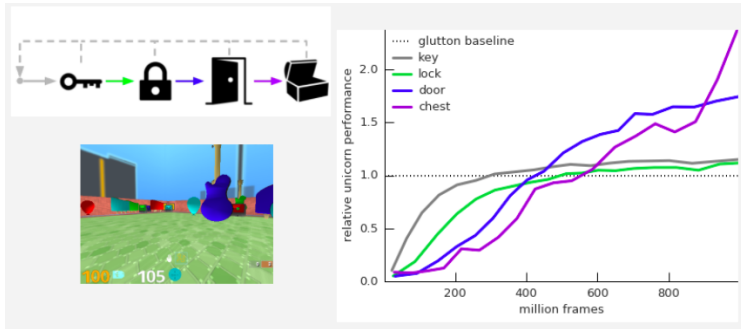
$$q_{c,\gamma,\pi}(s, a) = \mathbb{E}[C_{t+1} + \gamma_{t+1}C_{t+2} + \gamma_{t+1}\gamma_{t+2}C_{t+3} + \dots \mid S_t = s, A_{t+i} \sim \pi(S_{t+i})]$$

where  $C_t = c(S_t)$  and  $\gamma_t = \gamma(S_t)$  where  $S_t$  could be the environment state

- ▶  $c : \mathcal{S} \rightarrow \mathbb{R}$  is the **cumulant**
  - ▶ Predict many things, including — but not limited to — reward
- ▶  $\gamma : \mathcal{S} \rightarrow \mathbb{R}$  is the **discount** or termination
  - ▶ Predict for different time horizons  $\gamma$
- ▶  $\pi : \mathcal{S} \rightarrow \mathcal{A}$  is the **target policy**
  - ▶ Predict under many different (hypothetical) policies  $\pi$

# Universal value function approximators (Schaul et al. 2015)

- ▶ Idea: feed a representation of  $(c, \gamma)$  as input
- ▶ Allows generalization across goals/tasks within an environment



‘Unicorn’ (Mankowitz et al., 2018)

Learn about many things to learn to do the hard thing

## GVF and models

- ▶ A transition model is a specific set of GVFs
  - ▶ The cumulants are the state components (e.g., pixels)
  - ▶ The termination/discount is zero
  - ▶ Models are often action-conditional, so we do not need to specify the policy
- ▶ Similarly for an expected reward model
  - ▶ Cumulant = reward
  - ▶ Termination is immediate  $\gamma = 0$
- ▶ Easily extends to multi-step models
- ▶ Rolling forward a model = using predictions as inputs for other predictions

# Overview of the course

1. **Learning to make decisions** in bandit problems; exploration vs exploitation; learning action values; greedy and  $\epsilon$ -greedy; policy gradient for bandits; UCB
2. **Sequential decision problems**; MDPs; planning with dynamic programming; policy evaluation + policy improvement = policy iteration
3. **Model-free prediction and control**; Monte Carlo returns; TD learning; on-policy; off-policy; Q-learning; Sarsa; Double Q-learning
4. **Function approximation and deep RL**; tabular vs linear vs non-linear; convergence and divergence; least-squares prediction (LSTD and LSMC); multi-step returns; neural Q-learning; DQN
5. **Policy gradients and actor-critic methods**; REINFORCE; advantage actor-critics (A2C); trust-region methods; continuous actions; CACLA; gradient ascent on the value (DPG)
6. **Learning from a model**; Full models vs expectation models vs stochastic (generative) models; Dyna; parametric vs non-parametric models; experience replay; search; MCTS